

A Guide to Optimization and Optimal Play in Greed

December 14, 2023

Abstract

Abstract

The game of Greed is a probabilistic two-player game. This paper aims to derive an optimal way to play the game - based on reasonable heuristic metrics - as well as as determine the optimal chance of success, as determined by out heuristic.

In other word, this paper aims to make you the best greed player, so continue on and lord your superior greed strategy over your enemies.

How Do You Play Greed Again?

Game Environment & Turns

Greed is a game that could be described as “kinda like blackjack with dice”.

Game parameters:

M : max score before going bust.

s : number of sides on each die.

Starts at $s_0 = 0$ and $s_1 = 0$ respectively.

Each turn, the player up to roll will choose some $n \in \mathbb{N}^0$ number of dice to roll. Then the sum of those dice will be added to their score. Players will go back and forth like this until one of two terminal states is reached.

How To Win

1. In the first option, a player's score goes over M , i.e. and they go bust. In this case, they lose, so we'll say that their rating – the heuristic measure that we use to measure how good a position is – is 0, and thus their opponents rating is 1.
2. In the second option, a player can choose to roll 0 dice. This will initiate the last turn. The other player will then have one more chance to roll. The winner is the player who has the higher score that is not over M . In the case of a tie, rating is $1/2$ for each player.

Mathematical Framework

Mathematical Framework

Consider two types of states: terminal states, which denote the last turn of the game, and normal turns, representing all other states.

Normal and terminal states are conceptualized as 2 by 2 arrays with dimension $M + 1$ by $M + 1$, representing all possible states for player and opponent.

The x-axis designates the player presently up to roll, while the y-axis designates the player who has just rolled.

Therefore, any state S can be defined by the tuple (s_0, s_1, I)

PMF, PMF, PMF

What Is The Random Variable

First, consider just a single dice. It has pmf

$$D_i^{(s)}(d) = \begin{cases} \frac{1}{s} & \text{if } d \in \{1, \dots, s\} \\ 0 & \text{otherwise} \end{cases}$$

So let the random variable T denote the sum of n iid random dice, each with s sides. It can thusly be written

$$T : (\mathbb{N}^0)^n \rightarrow \mathbb{R}, T := D_1^{(s)}(d_1) + \dots + D_n^{(s)}(d_n)$$

Therefore, our goal is to find the pmf of T , which is notated $\mathbf{p}_T^{(n,s)}(t)$, dependent on parameters n, s .

Moment-Generating Functions

In order to find the probability mass function of T , we'll use moment generating functions. Remember that probability distributions and moment generating functions have a one-to-one correspondence. So for a discrete random variable X with probability distribution

$$f(x_i) = P(X = x_i) = p_i \text{ for } i = 1, 2, \dots, k$$

then its mfg is

$$\begin{aligned} M_X(t) &= \mathbb{E}[e^{tX}] = \sum_x e^{tx} \cdot f(x) = \\ &= p_1 \cdot e^{tx_1} + p_2 \cdot e^{tx_2} + \dots + p_k \cdot e^{tx_k}. \end{aligned}$$

and visa verse (like in [HW 6, 2.1]).

So for $t = 1$, $P(X = x)$ is the coefficient of e^x .

What is the moment generating function of T ?

Recall the pmf of $D_i^{(s)}$. It's moment generating function is defined as

$$M_{D_i^{(s)}}(t) = \mathbb{E}[e^{tD_i^{(s)}}] = \frac{1}{s}(e^t + e^{2t} + \dots + e^{st})$$

Since $D_i^{(s)}$ are all independent and identically distributed,

$$\begin{aligned} M_T(t) &= \mathbb{E}[e^{tT}] = \mathbb{E}[e^{t(D_1^{(s)} + \dots + D_n^{(s)})}] = \prod_{i=1}^n \mathbb{E}[e^{tD_i^{(s)}}] \\ &= \prod_{i=1}^n \left[\frac{1}{s}(e^t + e^{2t} + \dots + e^{st}) \right] \\ &= \frac{1}{s^n}(e^t + e^{2t} + \dots + e^{st})^n \end{aligned}$$

Formula for Coefficients of the Multinomials (Strap In)

Skipping over the complex algebra, the coefficient x^t in the expanded multinomial are found with

$$\frac{1}{s^n} \sum_{k=0}^{\lfloor \frac{t-n}{s} \rfloor} (-1)^k \binom{n}{k} \binom{t-sk-1}{n-1}$$

And as explained, these coefficients map to the probabilities of the pmf, so the pmf of T is written

$$\mathbf{p}_T^{(n,s)}(t) = \frac{1}{s^n} \sum_{k=0}^{\lfloor \frac{t-n}{s} \rfloor} (-1)^k \binom{n}{k} \binom{t-s \cdot k-1}{n-1}$$

Terminal States

Defining a Rating Function on Terminal States

The goal is to find some n to maximize the probability of getting a new score $s_p + t$ between s_o , and M ?"

Or more precisely, given a state $(s_p, s_o, 1)$, what is the optimal n such that the expected rating is maximized.

$$\text{rating}((s_p, s_o, 1), n) := \sum_{t=s_o+1}^M \mathbf{p}_T^{(n,s)}(t - s_p) + \frac{1}{2} \cdot \mathbf{p}_T^{(n,s)}(s_o - s_p)$$

where the summation describes the weighted sum of all next states given n , weighted according to their probability (transition matrix), hence rating is the expectation of its next possible states.

Optimal Actions and Ratings on Terminal States

Since n_\star is the optimizer of rating, it is given by

$$n_\star(s_0, s_1, 1) := \operatorname{argmax}_n \{\operatorname{rating}((s_0, s_1, 1), n)\}$$

Notice that the optimal rating comes for free. It's the rating that was optimized for in finding n_\star , so no additional work is required.

$$\operatorname{rating}_\star((s_p, s_o, 1), n_\star) := \operatorname{rating}((s_p, s_o, 1), n_\star)$$

Optimizing Normal States

Defining a Rating Function on Normal States

The rating is defined in the same way as it is for terminal states:
The expectation of the rating for the next possible states S_1 given some n .

Example

Let's give an example. Imagine that the optimal n_* and rating_* for every other state is known.

Consider rolling 2 dice at state $(2, 6)$. You could end up at any of the following states: $S_1 = \{(6, 4), (6, 5), (6, 6), (6, 7)\}$. So since we know the rating of all these states, we can calculate the rating given n by taking the weighted sum over S_1 , with weights given by the pmf of T , i.e. $\mathbf{p}_T^{(n,s)}(2), \dots, \mathbf{p}_T^{(n,s)}(4)$.

Defining a Rating Function on Normal States

Fact: Rating is complementary with respect to each player's rating for a given state.

Hence, the optimal rating for landing on a state S is $1 - \text{rating}(S, n_\star)$ (since the rating of that state will be for the opponent), where n_\star is the optimal n for that state.

So the rating function is given by

$$\text{rating}(S, n) := \begin{cases} \sum_{t=n}^{S \cdot n} 1 - \text{rating}((s_o, s_p + t, 0), n_\star) & \text{if } n > 0 \\ 1 - \text{rating}((s_o, s_p, 1), n_\star) & \text{if } n = 0 \end{cases}$$

Optimal Actions and Ratings on Normal States

Thus the optimal n_* given any possible state S , normal or terminal is now defined to be

$$n_* = \operatorname{argmax}_n \left\{ \sum_{t=n}^{s \cdot n} 1 - \operatorname{rating}((s_o, s_p + t, \{0 \text{ if } n = 1 \text{ else } 0\})) \right\}$$

where the rating function could be the one defined for terminal states or normal states.

rating_* is defined the same as for terminal states.

Cool Plots

Terminal States Results




Figure 1: Optimal Actions and Ratings for Terminal States

Normal States Results




Figure 2: Optimal Actions and Ratings for Normal States

Conclusion

Important Results

- Found a metric for determining success in Greed
- Found the pmf of n independent and identically distributed discrete uniform random variables, each with s outcomes.
- Found a rating function for all possible states that could occur in Greed.
- Found a optimal move and rating for every possible state in Greed

Limitations

- Rating is not an objective measurement
- The way that states rely on each other may not be optimal
- The code that runs the calculations could still have many bugs and possible edge cases (I would like to concur on the possibility of this outcome).

- A more mathematically rigorous mathematical framework.
- Determin whether there some pattern on normal states that can allow us to infer their rating more easily.
- etc

- Calculations for getting the coefficient of the multinomial
- Algorithm for finding the optimal n and rating for terminal states.
- Algorithm for finding the optimal n and rating for normal states.